Practical 7 : Sharding in MongoDB

1. Understanding the Components

A MongoDB **sharded cluster** consists of:

- **Config Servers (CSRS)** – Store metadata about the cluster.
- **Shards** – Hold the actual data.
- **Query Routers (mongos)** – Route client queries to the appropriate shard

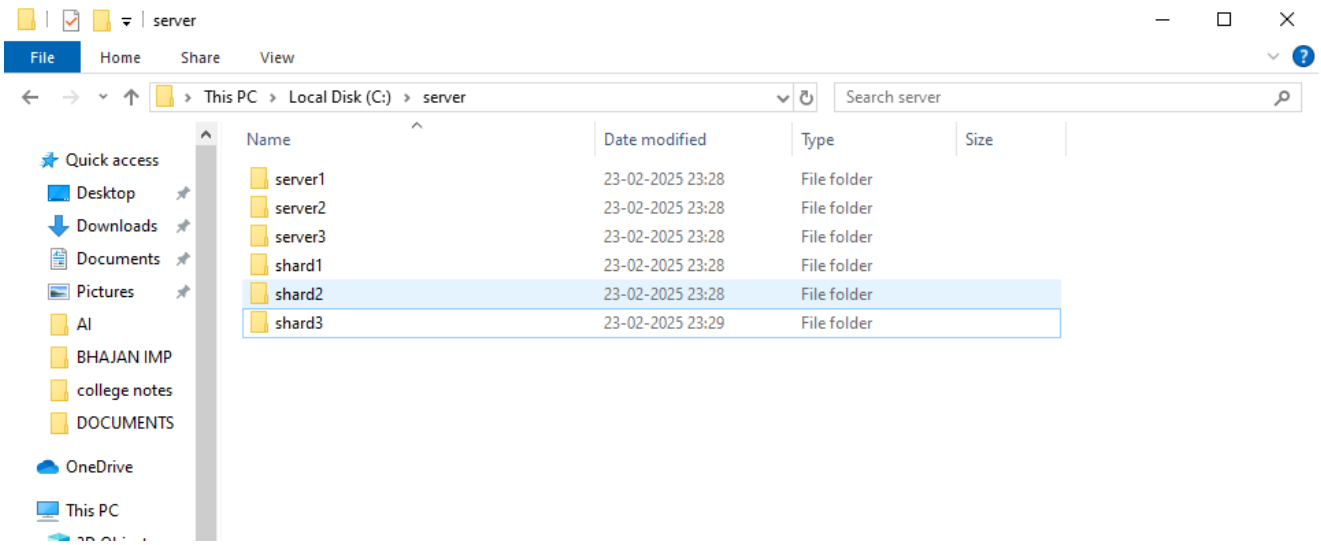

2. Setting Up Config Servers

Config servers form a **Replica Set** that maintains cluster metadata.

Where to Perform:

Run each command in a separate terminal window.

Start Config Servers

```
mongod --configsvr --port=1030 --replSet="configReplSet" --dbpath="/server/server1"
mongod --configsvr --port=1040 --replSet="configReplSet" --dbpath="/server/server2"
mongod --configsvr --port=1050 --replSet="configReplSet" --dbpath="/server/server3"
```

3. Initialize Config Server Replica Set

Connect to any one of the config servers using cmd:

Practical 7 : Sharding in MongoDB

```
Microsoft Windows [Version 10.0.19045.5487]
(c) Microsoft Corporation. All rights reserved.

C:\Users\Sejal>mongosh --host="localhost:1030"
Current Mongosh Log ID: 67bb5fca934245c4635d5f29
Connecting to:          mongodb://localhost:1030/?directConnection=true&serverSelectionTimeoutMS=2000&appName=mongosh+2.1.5
Using MongoDB:          7.0.7-rc0
Using Mongosh:          2.1.5
mongosh 2.4.0 is available for download: https://www.mongodb.com/try/download/shell

For mongosh info see: https://docs.mongodb.com/mongodb-shell/

------
   The server generated these startup warnings when booting
   2025-02-23T23:09:43.016+05:30: Access control is not enabled for the database. Read and write access to data and configuration is unrestricted
   2025-02-23T23:09:43.017+05:30: This server is bound to localhost. Remote systems will be unable to connect to this server. Start the server with --bind_ip <address>
to specify which IP addresses it should serve responses from, or with --bind_ip_all to bind to all interfaces. If this behavior is desired, start the server with --bind
_ip 127.0.0.1 to disable this warning
------
```

Initiate the replica set

```
test> rs.initiate({
...     _id: "configReplSet",
...     configsvr: true,
...     members: [
...         { _id: 0, host: "localhost:1030" },
...         { _id: 1, host: "localhost:1040" },
...         { _id: 2, host: "localhost:1050" }
...     ]
... })
{ ok: 1 }
configReplSet [direct: other] test>
```

4. Setting Up Shards

Shards store the actual data and are also configured as **Replica Sets**.

Where to Perform:

Run each command in a separate terminal window.

mongod --shardsvr --port=1130 --replSet="shardReplSet" --dbpath="/server/shard1"

mongod --shardsvr --port=1140 --replSet="shardReplSet" --dbpath="/server/shard2"

mongod --shardsvr --port=1150 --replSet="shardReplSet" --dbpath="/server/shard3"

```
Command Prompt
Microsoft Windows [Version 10.0.19045.5487]
(c) Microsoft Corporation. All rights reserved.

C:\Users\Sejal>mongod --shardsvr --port=1150 --replSet="shardReplSet" --dbpath="C:/server/shard3"
```

Practical 7 : Sharding in MongoDB

5. Initialize Shard Replica Set

Connect to any one of the shard servers using cmd

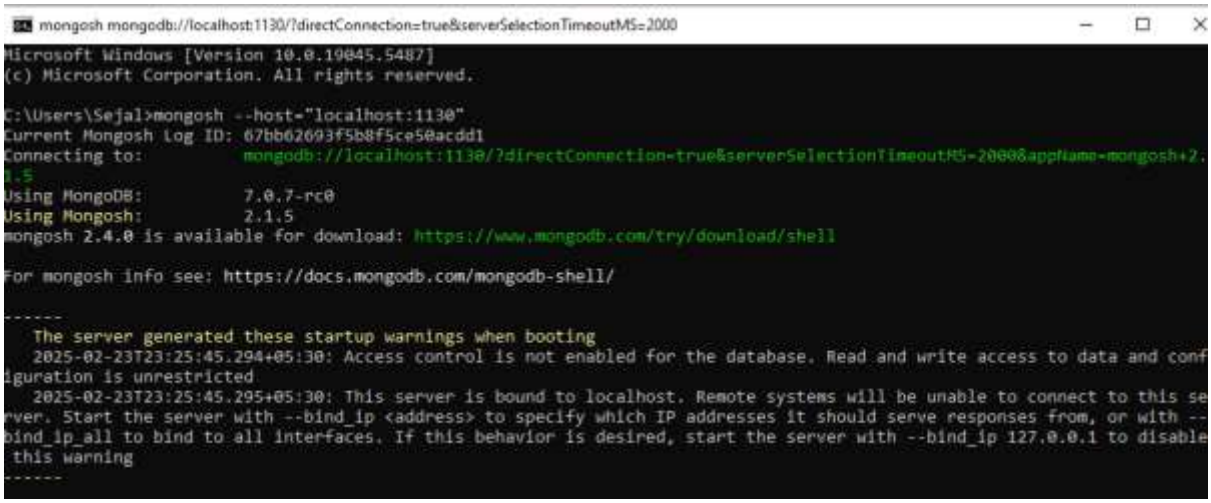

Inside cmd, initiate the replica set:

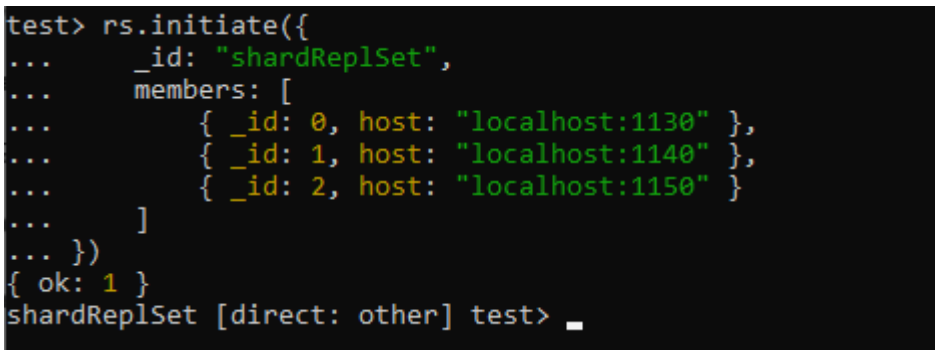

6. Setting Up Query Router (mongos)

The **mongos** process acts as a router that distributes queries across shards.

Where to Perform:

Run in a separate terminal.

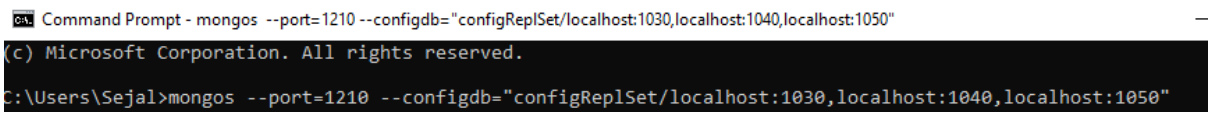

Practical 7 : Sharding in MongoDB

7. Connecting Shards to the Query Router

Connect to the query router using cmd:

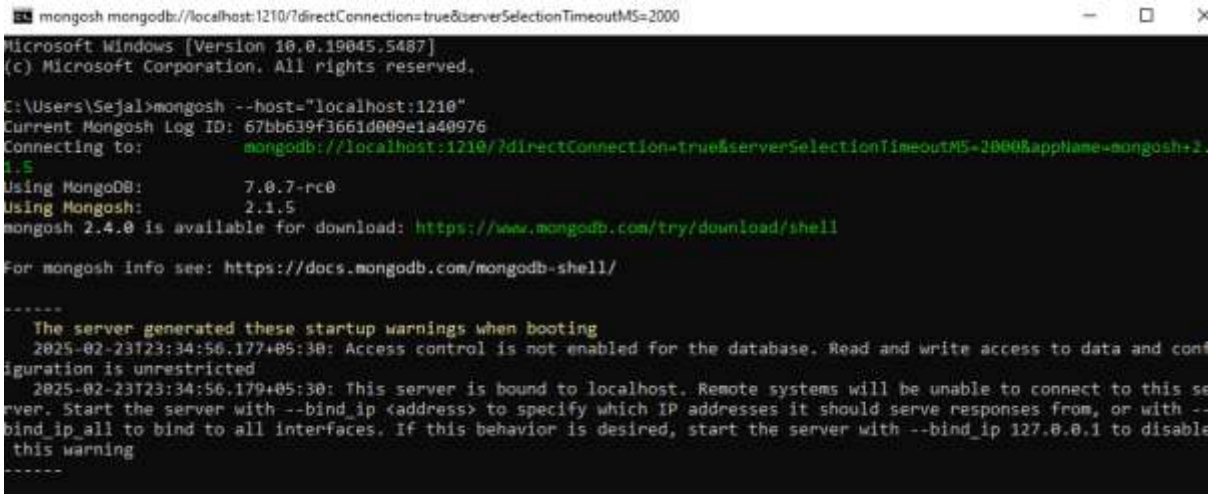

Inside cmd, add the shard replica set:

```
[direct: mongos] test> sh.addShard("shardReplSet/localhost:1130,localhost:1140,localhost:1150")
{
  shardAdded: 'shardReplSet',
  ok: 1,
  '$clusterTime': {
    clusterTime: Timestamp({ t: 1740334073, i: 4 }),
    signature: {
      hash: Binary.createFromBase64('AAAAAAAAAAAAAAAAAAAAAAAAAAA=', 0),
      keyId: Long('0')
    }
  },
  operationTime: Timestamp({ t: 1740334073, i: 4 })
}
[direct: mongos] test>
```

8. Enabling Sharding for a Database

Enable sharding for a specific database:

```
[direct: mongos] test> sh.enableSharding("practice")
{
  ok: 1,
  '$clusterTime': {
    clusterTime: Timestamp({ t: 1740334212, i: 1 }),
    signature: {
      hash: Binary.createFromBase64('AAAAAAAAAAAAAAAAAAAAAAAAAAA=', 0),
      keyId: Long('0')
    }
  },
  operationTime: Timestamp({ t: 1740334211, i: 3 })
}
[direct: mongos] test>
```

Practical 7 : Sharding in MongoDB

9. Sharding a Collection

Shard the students collection using enroll as the shard key:

```
[direct: mongos] test> sh.shardCollection("practice.students", { "enroll": 1 })
{
  collectionsharded: 'practice.students',
  ok: 1,
  '$clusterTime': {
    clusterTime: Timestamp({ t: 1740334412, i: 34 }),
    signature: {
      hash: Binary.createFromBase64('AAAAAAAAAAAAAAAAAAAAAAAAAAA=', 0),
      keyId: Long('0')
    }
  },
  operationTime: Timestamp({ t: 1740334412, i: 34 })
}
```

10. Testing the Sharded Cluster

To verify everything is working:

Practical 7 : Sharding in MongoDB

This will show details about shards, config servers, and sharded databases.

Insert data and check distribution:

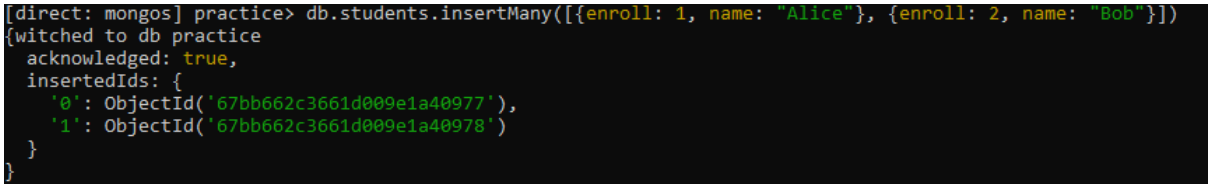

11. To confirm data exists on Shard2 (localhost:1140) and Shard3 (localhost:1150):

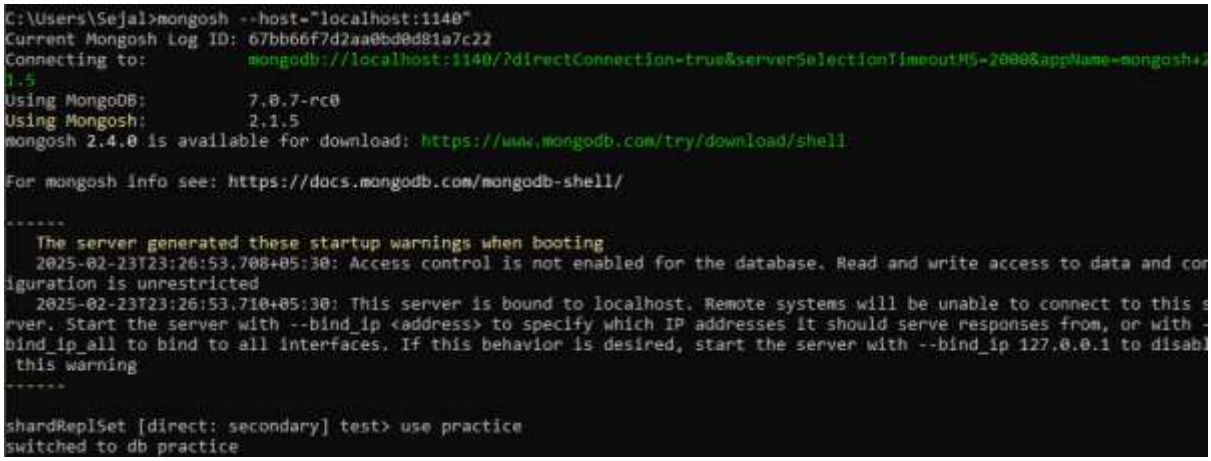

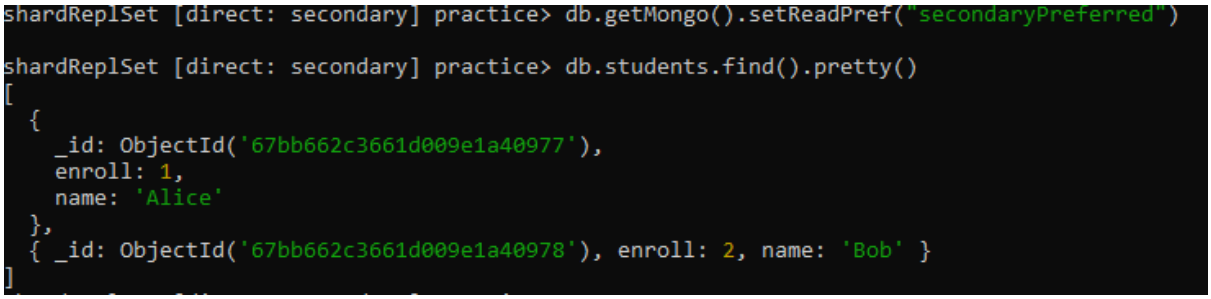

Practical 7 : Sharding in MongoDB